



Machine Learning

Reinforcement Learning (1)

Dr. Harry Goldingay
goldhj1@aston.ac.uk

Learning Outcomes

- ▶ At the end of this lecture you should:
 - ▶ Know what a sequential decision making problem is and how to formulate it as a Markov Decision Process (MDP).
 - ▶ Understand solution policies for MDPs and know a simple algorithm for generating an optimal policy given sufficient information.
 - ▶ Understand the characteristics of problems addressed by reinforcement learning and know an algorithm for policy evaluation on these problems well enough to implement it.

Motivation

Game Playing


- Some of the biggest machine learning media stories in recent years concern “game playing”.



- How would you apply machine learning to playing chess?

Learning Chess

- Firstly, what do we need to learn?



Rh1+

Learning Chess using Supervised Learning


- ▶ We want to learn a mapping from **chess positions** to **best moves**.
 - ▶ Chess positions could be seen as a **feature vector**.
 - ▶ Single, categorical output needed.
- ▶ Should remind you of **classification**.
 - ▶ We could apply supervised learning...
 - ▶ ...to a database of games from top players...
 - ▶ ...to try to predict how they would move in the position.
- ▶ Issues:
 - ▶ How can we improve on top players by copying them?
 - ▶ Do we have enough (of the right) data?

Sequential Decision Making

Sequential Decision Making

- ▶ A more useful way to think of this problem is as one of **sequential decision making**:
 - ▶ an agent takes an **action**,
 - ▶ it **observes** the effect of the action...
 - ▶ ...and then chooses a subsequent action.
- ▶ This generates a sequence of actions and effects.
- ▶ The aim is to take actions which lead to the “best” sequence of effects.

Grid World




“Best” Sequence

- ▶ On each square in grid world, the agent will receive a **reward**:
 - ▶ +1 on square (4,3)
 - ▶ -1 on square (4,2)
 - ▶ -0.04 on any other square.
- ▶ If the agent reaches either of squares (4,3) or (4,2) (**terminal states**), the game ends.
- ▶ The best sequence will be the one which maximises **utility**: the sum of rewards.
 - ▶ Note: there are other ways of defining utility.

Solving Grid World

- ▶ A “solution” is just a sequence of actions.
- ▶ The **rewards** an agent receives in grid world...
 - ▶ +1 or -1 for a transition to an exit (terminal state),
 - ▶ -0.04 for a transition to any other square,
- ▶ ...mean that it should find its way to the terminal at (4,3) as quickly as possible.
- ▶ The optimal solution is the sequence which achieves the above.


Grid World: Optimal Solution(s)



Decision Making with Uncertainty

- ▶ That was easy but, unfortunately, unrealistic.
 - ▶ Many interesting decision problems involve uncertainty.
 - ▶ Our actions effect outcomes, NOT determine them.
- ▶ We can incorporate uncertainty into our problem, by making transitions probabilistic:
 - ▶ Make the intended move with probability 0.8,
 - ▶ Make a perpendicular move with probability 0.2 (each direction equally likely),
 - ▶ If the agent hits a wall, it stays where it is.
- ▶ How does this affect our solutions?


Stochastic Grid World: Optimal Solutions?



Policies

- ▶ In a stochastic setting, we prefer the path **Up, Up, Right, Right, Right** to the other optimal deterministic path.
- ▶ Is this a good solution to the problem?
 - ▶ No!
 - ▶ What if one of the steps fails (e.g. we go Right instead of Up) at step 1?


Stochastic Grid World: Optimal Solutions?



Policies

- ▶ In a stochastic setting, we prefer the path **Up, Up, Right, Right, Right** to the other optimal deterministic path.
- ▶ Is this a good solution to the problem?
 - ▶ No!
 - ▶ What if one of the steps fails (e.g. we go Right instead of Up) at step 1?
- ▶ A sequence of actions is not a good solution if we don't know the results of the actions.
 - ▶ Better to view each location as a **state**...
 - ▶ ...and formulate a solution by choosing an action per state.
 - ▶ Such a solution is called a **policy**.


Stochastic Grid World: Optimal Policy




Markov Decision Processes

- ▶ Sequential decision problems are often formalised as **Markov Decision Processes (MDPs)** which require:
 - ▶ A set of **states**, S , including a start state, s_0
 - ▶ For each state, a set of **actions**, A_s
 - ▶ A **transition model**, $P(s'|s, a)$
 - ▶ Gives the probability of reaching state s' from state s if taking action a .
 - ▶ A **reward function**, $R(s)$
 - ▶ Sometimes also dependent on action and outcome:
 $R(s, a, s')$
- ▶ The goal is to find an optimal policy, π^* , which maximises **expected utility**.


Rewards and Optimal Policy




$$R(s) < -1.6284$$



$$-0.4278 < R(s) < -0.0850$$




$$-0.0221 < R(s) < 0$$



$$R(s) > 0$$


MDPs: Expected Utility

Playing an MDP (1)




$$\begin{aligned} R(s_0) &= -0.50 \\ R(s_1) &= -0.75 \\ R(s_2) &= -0.10 \end{aligned}$$

Playing an MDP (2)



$$\begin{aligned} R(s_0) &= 0.50 \\ R(s_1) &= 0.25 \\ R(s_2) &= 0.90 \end{aligned}$$

Playing an MDP (3)



$$\begin{aligned} R(s_0) &= -0.50 \\ R(s_1) &= -0.75 \end{aligned}$$

Utilities Over Time

- ▶ So far, we have been assigning utility on the basis of **additive rewards**:

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$$

- ▶ It is common to use discounted rewards:

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

- ▶ Where $0 < \gamma < 1$ is called a **discount factor**. Advantages:
 - ▶ Utility converges even for an infinite sequence of actions,
 - ▶ Seems to be a good model for human preferences.

Utilities of States

- ▶ We can define the utility of a state under policy π as:

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$

- ▶ where $S_0 = s$.

- ▶ The optimal policy for a given state is:

$$\pi_s^* = \operatorname{argmax}_{\pi} U^\pi(S)$$

- ▶ We can show that, for any state, $\pi_s^* = \pi^*$. Therefore, we can denote the (expected) utility of a state as:

$$U(s) = U^{\pi^*}(s)$$

Grid World: Utilities of States

3	0.812	0.868	0.918	+ 1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388
	1	2	3	4

For $\gamma = 1$

Policies and Utilities of States

- ▶ If we know $U(s)$ for each state, we can use it to infer the optimal policy.
 - ▶ Choose the action which results in the highest expected utility of the next state.
- ▶ What action should an agent take in state (3,1) in the previous example?

Grid World: Utilities of States

3	0.812	0.868	0.918	+ 1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388
	1	2	3	4

For $\gamma = 1$

Policies and Utilities of States

- ▶ If we know $U(s)$ for each state, we can use it to infer the optimal policy.
 - ▶ Choose the action which results in the highest expected utility of the next state.
- ▶ What action should an agent take in state (3,1) in the previous example?
 - ▶ Up: $0.8 \times 0.66 + 0.1 \times (0.655 + 0.388) \approx 0.63$
 - ▶ Left: $0.8 \times 0.655 + 0.1 \times (0.66 + 0.611) \approx 0.65$
- ▶ The agent should choose Left.
- ▶ **Key takeaway:** π^* can be inferred directly from utility.

Calculating Expected Utility

- ▶ Does this give us a solution methodology for MDPs?
- ▶ Not yet...
 - ▶ We can calculate π^* from $U(s)$
 - ▶ ...but recall that $U(s) = U^{\pi^*}(s)$
- ▶ ...but there iterative are algorithms for doing so:
 - ▶ Value iteration
 - ▶ Policy iteration

Value Iteration

- ▶ A consequence of our definition of utility is that the utility of states can be expressed as follows (**Bellman equation**):

$$U(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s'|s, a) U(s')$$

- ▶ This gives us a system of equations to solve (one for each state) but the max component makes this problematic.
- ▶ Try an iterative approach.

Value Iteration


- ▶ We can turn the Bellman equations into update equations:

$$U_{i+1}(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s'|s, a) U_i(s')$$

- ▶ If $\gamma < 1$, then this is guaranteed to converge to the (unique) solution to the Bellman equations.
- ▶ In practice:
 - ▶ Pick arbitrary initial values for the $U_0(s)$
 - ▶ Run until the difference between estimates at different timesteps become small

▶ This is called **value iteration**.

Value Iteration Illustrated



Value Iteration Illustrated

Transition Model

$$\begin{aligned}P(s_0|s_0, a_0) &= 0.5 \\P(s_2|s_0, a_0) &= 0.5 \\P(s_1|s_0, a_1) &= 1.0 \\P(s_0|s_1, a_0) &= 1.0 \\P(s_1|s_1, a_1) &= 0.1 \\P(s_2|s_1, a_1) &= 0.9\end{aligned}$$

Reward Function

$$\begin{aligned}R(s_0) &= -0.50 \\R(s_1) &= -0.75 \\R(s_2) &= -0.10\end{aligned}$$

Utility Estimates ($i = 0$)

$$\begin{aligned}U_i(s_0) &= 0 \\U_i(s_1) &= 0 \\U_i(s_2) &= 0\end{aligned}$$

Update

Initial choice of action unimportant. All states have estimated utility 0.

$$\begin{aligned}U_{i+1}(s_0) &= R(s_0) + 0 = -0.50 \\U_{i+1}(s_1) &= R(s_1) + 0 = -0.75 \\U_{i+1}(s_2) &= R(s_2) + 0 = -0.10\end{aligned}$$

Value Iteration Illustrated

Transition Model

$$\begin{aligned}P(s_0|s_0, a_0) &= 0.5 \\P(s_2|s_0, a_0) &= 0.5 \\P(s_1|s_0, a_1) &= 1.0 \\P(s_0|s_1, a_0) &= 1.0 \\P(s_1|s_1, a_1) &= 0.1 \\P(s_2|s_1, a_1) &= 0.9\end{aligned}$$

Reward Function

$$\begin{aligned}R(s_0) &= -0.50 \\R(s_1) &= -0.75 \\R(s_2) &= -0.10\end{aligned}$$

Utility Estimates ($i = 1$)

$$\begin{aligned}U_i(s_0) &= -0.50 \\U_i(s_1) &= -0.75 \\U_i(s_2) &= -0.10\end{aligned}$$

Update Example

Expected value of next state from s_0 given a .

$$a_0: 0.5 \times -0.5 + 0.5 \times -0.1 = -0.3$$

$$a_1: 1.0 \times -0.75 = -0.75$$

Best action: a_0 .

$$U_{i+1}(s_0) = R(s_0) - 0.3 = -0.8$$

Value Iteration Illustrated

Transition Model

$$\begin{aligned}P(s_0|s_0, a_0) &= 0.5 \\P(s_2|s_0, a_0) &= 0.5 \\P(s_1|s_0, a_1) &= 1.0 \\P(s_0|s_1, a_0) &= 1.0 \\P(s_1|s_1, a_1) &= 0.1 \\P(s_2|s_1, a_1) &= 0.9\end{aligned}$$

Reward Function

$$\begin{aligned}R(s_0) &= -0.50 \\R(s_1) &= -0.75 \\R(s_2) &= -0.10\end{aligned}$$

Utility Estimates ($i = 2$)

$$\begin{aligned}U_i(s_0) &= -0.8 \\U_i(s_1) &= -0.915 \\U_i(s_2) &= -0.10\end{aligned}$$

Update Example

Expected value of next state from s_0 given a .

$$a_0: 0.5 \times -0.8 + 0.5 \times -0.1 = -0.45$$

$$a_1: 1.0 \times -0.915 = -0.915$$

Best action: a_0 .

$$U_{i+1}(s_0) = R(s_0) - 0.45 = -0.95$$

Value Iteration Illustrated

Transition Model

$$\begin{aligned}P(s_0|s_0, a_0) &= 0.5 \\P(s_2|s_0, a_0) &= 0.5 \\P(s_1|s_0, a_1) &= 1.0 \\P(s_0|s_1, a_0) &= 1.0 \\P(s_1|s_1, a_1) &= 0.1 \\P(s_2|s_1, a_1) &= 0.9\end{aligned}$$

Reward Function

$$\begin{aligned}R(s_0) &= -0.50 \\R(s_1) &= -0.75 \\R(s_2) &= -0.10\end{aligned}$$

Utility Estimates ($i = 3$)

$$\begin{aligned}U_i(s_0) &\approx -0.95 \\U_i(s_1) &\approx -0.93 \\U_i(s_2) &\approx -0.10\end{aligned}$$

Utility Estimates ($i = 4$)

$$\begin{aligned}U_i(s_0) &\approx -1.03 \\U_i(s_1) &\approx -0.93 \\U_i(s_2) &\approx -0.10\end{aligned}$$

Utility Estimates ($i = 5$)

$$\begin{aligned}U_i(s_0) &\approx -1.06 \\U_i(s_1) &\approx -0.93 \\U_i(s_2) &\approx -0.10\end{aligned}$$

Change < 0.05

Introduction to Reinforcement Learning

Planning and Reinforcement Learning

- ▶ So far, the problems we have seen can be solved through **offline planning**:
 - ▶ Consider the information available...
 - ▶ ...and infer the optimal policy.
 - ▶ No need to take any actions!
- ▶ A reinforcement learning problem is an MDP in which we don't know:
 - ▶ the **transition model**,
 - ▶ The **reward function**.
- ▶ Not enough information to solve offline. Need to experiment!

Naïve Approach

- ▶ We know how to solve MDPs when we do know the transition model and reward function.
 - ▶ Try to learn them from the environment!
 - ▶ Pick policy/policies.
 - ▶ Count state transitions that we observe and average to infer transition probabilities.
 - ▶ Observe rewards directly.
- ▶ Assume that our model is correct, and solve using value iteration.
- ▶ How does this scale with the number of actions/states?


Types of Reinforcement Learning

- ▶ Solving an MDP without prior knowledge of transitions and rewards is challenging.
- ▶ We will leave it to one side for now and look at approaches to **policy evaluation**.
 - ▶ Given a policy, what are the expected utilities of the states?
- ▶ This type of approach, with a fixed policy, is called **passive learning**.
- ▶ This is in contrast to **active learning**, in which an agent must also decide which actions to take.

Direct Evaluation

- ▶ We can take a similar approach to policy evaluation to our initial suggestion for solving RL problems:
- ▶ Repeatedly run trials of a policy from random starting states.
- ▶ Whenever we encounter a state, record the utility which follows in the trial.
- ▶ Average these records, to estimate the utility of the state under the policy.

Direct Evaluation Illustrated



$$\begin{aligned} R(s_0) &= -0.50 \\ R(s_1) &= -0.75 \\ R(s_2) &= -0.10 \end{aligned}$$

Direct Evaluation Illustrated

- ▶ Running four random trials with the optimal policy gives me:
 - ▶ $s_2(-0.1)$
 - ▶ $s_0(-0.5) \rightarrow s_2(-0.1)$
 - ▶ $s_1(-0.75) \rightarrow s_2(-0.1)$
 - ▶ $s_0(-0.5) \rightarrow s_0(-0.5) \rightarrow s_0(-0.5) \rightarrow s_2(-0.1)$
- ▶ We have seen s_0 four times, with an average utility of -0.975 :
$$\frac{-0.6 - 1.6 - 1.1 - 0.6}{4}$$
- ▶ We have seen s_1 once, with a total utility of -0.85
- ▶ We have seen s_2 four times, with an average utility of -0.1

Direct Evaluation

- ▶ Direct evaluation will eventually converge to the correct utility values.
- ▶ We don't need to learn the full transition model and reward function
 - ▶ don't have the same problem with number of state-action pairs as we saw before.
- ▶ However, the Bellman equations give us relationships between states.
 - ▶ Direct evaluation doesn't make use of these relationships...
 - ▶ ...so is very inefficient.

Temporal Difference Learning

Temporal Difference Learning

- ▶ Rather than trying to build a model of transitions directly, **temporal difference learning** operates directly on relationships between states.
- ▶ Whenever an agent transitions from state s to state s' (in which it received rewards $R(s)$ and $R(s')$ respectively), it updates its estimate of utility as follows:

$$U(s') = R(s') \quad \text{if } s' \text{ has not been visited previously}$$

$$U(s) = U(s) + \alpha(N_s(s))(R(s) + \gamma U(s') - U(s))$$

- ▶ where $\alpha(N_s(s))$ is a “step-size”, decreasing in the number of times the state s is encountered (e.g. $\frac{1}{1+N_s(s)}$).


TDL Breakdown

- Breaking down the equation:

$$U(s) + \alpha(N_s(s)) \underbrace{(R(s) + \gamma U(s') - U(s))}_{\text{Estimate of utility of } s \text{ based on transition in current episode} \quad \text{Previous estimate of utility of } s},$$

- This difference $(R(s) + \gamma U(s') - U(s))$ will be positive if the estimate of utility based on the current episode is higher than our previous estimate.
- A positive difference would be evidence that our estimate of utility $(U(s))$ was too low and should be increased.
- $\alpha(N_s(s))$ decreases as we encounter state s more. Each new observation is a smaller fraction of our total observations.

Temporal Difference Learning Illustrated



Temporal Difference Learning Illustrated

- ▶ Imagine that the previous MDP was encountered twice by an agent whose policy was to play a_1 in both states, generating the following sequences of states and rewards:
- ▶ $s_0(-0.5) \rightarrow s_1(-0.75) \rightarrow s_2(-0.1)$
- ▶ $s_1(-0.75) \rightarrow s_1(-0.75) \rightarrow s_2(-0.1)$
- ▶ How to apply TDL?

Temporal Difference Learning Illustrated

Transitions

→ $s_0(-0.5)$
 $s_0(-0.5) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

 $s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

Counts

$$N_S(s_0) = 0$$
$$N_S(s_1) = 0$$
$$N_S(s_2) = 0$$

Utility Estimates

$$U(s_0) = ?$$
$$U(s_1) = ?$$
$$U(s_2) = ?$$

Update

State s_0 encountered for the first time.

$$U(s_0) = -0.5$$

Temporal Difference Learning Illustrated

Transitions

$s_0(-0.5)$
 $s_0(-0.5) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

 $s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

Counts

$N_s(s_0) = 0$
 $N_s(s_1) = 0$
 $N_s(s_2) = 0$

Utility Estimates

$U(s_0) = -0.5$
 $U(s_1) = ?$
 $U(s_2) = ?$

Update

State s_1 encountered for the first time.

$$U(s_1) = -0.75$$

State s_0 updated based on observed transition to s_1 .

$$N_s(s_0) = 1$$

$$U(s_0) = -0.5 + \frac{1}{1+1} (-0.5 - 0.75 + 0.5) = -0.875$$

Diagram illustrating the update formula for $U(s_0)$ with red arrows pointing to the terms:

- $U(s)$ points to -0.5
- $N_s(s)$ points to the denominator $1+1$
- $R(s)$ points to the first -0.5 in the numerator
- $U(s')$ points to -0.75 in the numerator
- $U(s)$ points to the second 0.5 in the numerator

Temporal Difference Learning Illustrated

Transitions

$s_0(-0.5)$
 $s_0(-0.5) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$
 $s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

Counts

$N_s(s_0) = 1$
 $N_s(s_1) = 0$
 $N_s(s_2) = 0$

Utility Estimates

$U(s_0) = -0.875$
 $U(s_1) = -0.75$
 $U(s_2) = ?$

Update

State s_2 encountered for the first time.

$$U(s_2) = -0.1$$

State s_1 updated based on observed transition to s_2 .

$$N_s(s_1) = 1$$


$$U(s_1) = -0.75 + \frac{1}{1+1} (-0.75 - 0.1 + 0.75) = -0.8$$

Diagram annotations:

- $U(s)$ (red arrow pointing to -0.75)
- $N_s(s)$ (red arrow pointing to 1)
- $R(s)$ (red arrow pointing to -0.1)
- $U(s')$ (red arrow pointing to 0.75)
- $U(s)$ (red arrow pointing to the final result -0.8)

Temporal Difference Learning Illustrated

Transitions

$s_0(-0.5)$
 $s_0(-0.5) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$
 $s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

Counts

$N_S(s_0) = 1$
 $N_S(s_1) = 1$
 $N_S(s_2) = 0$

Utility Estimates

$U(s_0) = -0.875$
 $U(s_1) = -0.8$
 $U(s_2) = -0.1$

Update

State s_1 encountered previously. No change.

Temporal Difference Learning Illustrated

Transitions

$s_0(-0.5)$
 $s_0(-0.5) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

 $s_1(-0.75)$
→ $s_1(-0.75) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

Counts

$N_s(s_0) = 1$
 $N_s(s_1) = 1$
 $N_s(s_2) = 0$

Utility Estimates

$U(s_0) = -0.875$
 $U(s_1) = -0.8$
 $U(s_2) = -0.1$

Update

State s_1 encountered previously. No change.

State s_1 updated based on observed transition to s_1 .

$N_s(s_1) = 2$

$$U(s_1) = -0.8 + \frac{1}{1+2} (-0.75 - 0.8 + 0.8) = -1.05$$


Diagram illustrating the update formula for $U(s_1)$ based on the observed transition to s_1 . Red arrows point to the components of the formula:

- $U(s)$ points to -0.8
- $N_s(s)$ points to $1+2$
- $R(s)$ points to -0.75
- $U(s')$ points to 0.8

Temporal Difference Learning Illustrated

Transitions

$s_0(-0.5)$
 $s_0(-0.5) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

 $s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_1(-0.75)$
 $s_1(-0.75) \rightarrow s_2(-0.1)$

Counts

$N_s(s_0) = 1$
 $N_s(s_1) = 2$
 $N_s(s_2) = 0$

Utility Estimates

$U(s_0) = -0.875$
 $U(s_1) = -1.05$
 $U(s_2) = -0.1$






Update

State s_1 encountered previously. No change.

State s_1 updated based on observed transition to s_2 .

$N_s(s_1) = 3$

$$U(s_1) = -1.05 + \frac{1}{1+3} (-0.75 - 0.1 + 1.05) = -1$$

 $U(s)$  $N_s(s)$  $R(s)$  $U(s')$  $U(s)$

Conclusion

- ▶ You should know:
 - ▶ Know what a sequential decision making problem is and how to formulate it as a Markov Decision Process (MDP).
 - ▶ Understand solution policies for MDPs and know a simple algorithm for generating an optimal policy given sufficient information.
 - ▶ Understand the characteristics of problems addressed by reinforcement learning and know an algorithm for policy evaluation on these problems well enough to implement it.